

# MEILING WANG

West Lafayette, IN, USA 47906 ◊ wang7395@purdue.edu

## EDUCATION

---

**Purdue University**

B.S. in Computer Science

*Expected Graduation: Dec 2027*

GPA: 4.0/4.0

**Relevant Coursework:**

- **Graduate Level:** Artificial Intelligence, ML System Co-Design, Compilers
- **Undergraduate Level:** Data Structures, Analysis of Algorithms, Computer Architecture, Game Theory, NLP, Discrete Mathematics, Complex Analysis

## RESEARCH EXPERIENCE

---

**EcoAI Lab**, Research Assistant

*Purdue University*

*Mar 2026 – Present*

*Advisor: Prof. Haoran You*

- Co-authored a paper on efficient inference for diffusion large language models (dLLMs), under review at NeurIPS 2026.
- Designed and evaluated inference acceleration methods for dLLMs through large-scale experiments on multi-GPU (8×H100) infrastructure.

**EXcel Lab**, Research Intern

*May 2026 – Present*

*Advisor: Prof. Yuzhang Shang*

- Leading the development of a multi-agent system for autonomously authoring and optimizing multi-GPU communication CUDA kernels, evaluated on multi-node H100 clusters.

## PROJECTS

---

**Adversarial Robustness of Visual In-Context Learning**

*PurSec Lab*

*Aug 2025 – Dec 2025*

*Advisor: Prof. Antonio Bianchi*

- Built an evaluation pipeline for Visual In-Context Learning robustness under adversarial perturbations, measuring cross-task hijacking behavior across vision tasks.
- Developed task-level hijacking attacks that redirect model inference between vision tasks, achieving 81.7% task-switching success and 90% success on a real-world labeling pipeline.

**Trustworthy Machine Learning**

*Purdue VIP Program*

*Jan 2026 – May 2026*

*Advisor: Prof. Zahra Ghodsi*

- Deployed and benchmarked *Landseer* on Purdue's Gilbreth HPC cluster, enabling GPU-accelerated evaluation of privacy-preserving ML inference and backdoor defenses for fine-tuned models.

**Diffusion LLM Serving Framework**

*To be open-sourced*

*Jun 2026 – Present*

- Contributing model support to serving framework for diffusion language models (release in preparation), integrating new dLLM architectures into the inference stack with parallel-decoding and KV-cache-style optimizations.

## SKILLS & HONORS

---

**Programming Languages:** Python, Java, C, C++, Bash, MATLAB

**Frameworks & Tools:** PyTorch, Hugging Face Transformers, vLLM, Slurm, Git

**Teaching:** Teaching Assistant, ECE 57000 Artificial Intelligence (Graduate Level)